

Analysis of Algorithms - Assignment 3

Zhenrui Zheng

Chinese University of Hong Kong, Shenzhen
225040512@link.cuhk.edu.cn

Contents

Amortized Analysis	1
1.1 Queue Implementation with Two Stacks	1
1.2 Amortized Computation Cost Analysis	2
Element Selection Algorithm	3
2.1 Randomized Algorithm	3
2.2 Success Probability Analysis	3
Shelf Scheduling Algorithm	4
3.1 NP-Hardness Proof	4
3.2 3-Approximation Algorithm	4
3.3 Lower Bounds for Optimal Makespan	5
3.4 3-Approximation Proof	5
Randomized Algorithm	6

Amortized Analysis

1.1 Queue Implementation with Two Stacks

Implementation

We implement a queue using two stacks by maintaining:

- **Input Stack (S_{in}):** For enqueue operations
- **Output Stack (S_{out}):** For dequeue operations

Algorithm 1 Queue Operations with Two Stacks

Require: Two stacks S_{in} and S_{out}

```
1: function Push( $x$ )
2:    $S_{in}.\text{push}(x)$ 
3: end function
4: function Get()
5:   if  $S_{out}$  is empty then
6:     while  $S_{in}$  is not empty do
7:        $S_{out}.\text{push}(S_{in}.\text{pop}())$ 
8:     end while
9:   end if
10:  return  $S_{out}.\text{pop}()$ 
11: end function
```

Proof of O(1) Amortized Time

We use the *accounting method* to prove the amortized cost analysis.

Assigning Costs: We assign an amortized cost of 4 for each Push operation:

- 1 unit pays for the actual Push into S_{in}
- 3 unit is stored as credit on the element for its future operations

For Get operations, we assign an amortized cost of 0:

- When S_{out} is non-empty: the actual cost is 1 (one Pop), and we use 1 credit from the popped element (in S_{out})
- When S_{out} is empty: we transfer all elements from S_{in} to S_{out}
 - Each transfer requires one Pop from S_{in} (cost 1) and one Push to S_{out} (cost 1), paid by the credit on the element
 - After transfer, each element in S_{out} has 1 credit left, which is enough to cover the cost of the future Pops. Then, we can simply Pop 1 element from S_{out} with cost 1, covered by the credit on the element
 - Total cost: $2k + 1$ for transferring k elements and one final Pop, covered by $2k + 1$ credits

Invariant: At any point, the total credit stored equals $3|S_{in}| + |S_{out}|$, since each element in S_{in} has 3 units of credit from its Push operation, and each element in S_{out} has 1 credit left after transferring from S_{in} , which cost 2 credits.

Conclusion: Each Push has amortized cost $O(1)$, each Get has amortized cost $O(1)$, and the total amortized cost over any sequence of n operations is $O(n)$, giving $O(1)$ amortized cost per operation.

1.2 Amortized Computation Cost Analysis

Total Cost Calculation:

Let $k = \lfloor \log_2(n) \rfloor$. The total cost is:

$$\begin{aligned} T(n) &= \sum_{j=0}^k 2^j + (n - (k+1)) \cdot 1 \\ &= \sum_{j=0}^k 2^j + n - k - 1 \end{aligned}$$

We know that $\sum_{j=0}^k 2^j = 2^{k+1} - 1 = 2 \cdot 2^k - 1$.

Since $2^k \leq n < 2^{k+1}$, we have:

$$\begin{aligned} 2^k &\leq n \\ 2^{k+1} &\leq 2n \end{aligned}$$

Therefore:

$$\begin{aligned} T(n) &= 2^{k+1} - 1 + n - k - 1 \\ &\leq 2n - 1 + n - k - 1 \\ &= 3n - k - 2 \\ &\leq 3n \\ \frac{T(n)}{n} &\leq \frac{3n}{n} = 3 = O(1) \end{aligned}$$

Thus, the amortized computation cost per day is $O(1)$.

Element Selection Algorithm

2.1 Randomized Algorithm

Algorithm Description

Algorithm 2 Simplified Randomized Search

Require: Arrays $X[1..n]$, $next[1..n]$, $c > 0$ and value x

Ensure: True if $x \in X$, False otherwise

```
1: Let  $m = c\lceil\sqrt{n}\rceil$ 
2: Randomly sample  $m$  distinct indices uniformly from  $[1..n]$ 
3: Among the sampled indices, find  $i^*$  that minimizes  $|X[i] - x|$ 
4: Follow  $next$  pointers from  $i^*$  for at most  $m$  steps
5: for  $k = 1$  to  $m$  do
6:   if  $X[i^*] = x$  then
7:     return True
8:   end if
9:    $i^* = next[i^*]$ 
10: end for
11: return False
```

This requires $O(\sqrt{n})$ time complexity as expected.

2.2 Success Probability Analysis

Let X' be the sorted version of X , and x be the target element. Let I be the set of indices within distance $m = c\sqrt{n}$ from x in X' , that is, $I = \{i \in [1..n] \mid lowerbound(x) - i \leq m \pmod n\}$. The probability that, if we take m samples uniformly from $[1..n]$, at least one of them is in I is a hypergeometric distribution, which is given by:

$$\begin{aligned} P(\text{at least one sample among } m \text{ in } I) \\ = 1 - P(\text{all } m \text{ samples in } \bar{I}) \\ = 1 - \frac{\binom{m}{0} \binom{n-m}{m}}{\binom{n}{m}} \\ = 1 - \frac{(n-m)!(n-m)!}{n!(n-2m)!} \quad (m! \text{ cancels out)} \end{aligned}$$

We want this to be at least 0.99:

$$\begin{aligned} 1 - \frac{(n-m)!(n-m)!}{n!(n-2m)!} &\geq 0.99 \\ \frac{(n-m)!(n-m)!}{n!(n-2m)!} &\leq 0.01 \\ m &\leq 2.15\sqrt{n} \end{aligned}$$

So $m = \lceil 2.15\sqrt{n} \rceil$ suffices for 99% success probability. Similarly, minimum m for 99.9% and 99.99% success probabilities are $\lceil 2.63\sqrt{n} \rceil$ and $\lceil 3.04\sqrt{n} \rceil$, respectively.

Shelf Scheduling Algorithm

3.1 NP-Hardness Proof

We prove that the parallel scheduling problem is NP-hard by reduction from the Subset-sum problem under a special case, where $p = 2$.

Reduction: Given an instance of subset-sum with a set of positive integers $\{a_1, a_2, \dots, a_n\}$ and target sum S , we construct a 2-processor parallel scheduling instance as follows:

- occupy time: $T[1, \dots, n] = \{a_1, a_2, \dots, a_n\}$
- Processor: $P[1, \dots, n], P[i] = 1$

The question becomes: Can we schedule these n jobs on 2 processors such that all jobs finish at time $t = S$? This is equivalent to asking whether we can partition the n elements into two subsets, with one of them sum up to S .

The reduction is given within polynomial time, and since Subset-sum is NP-complete, the parallel scheduling problem is NP-hard.

3.2 3-Approximation Algorithm

We present a scheduling algorithm that achieves a 3-approximation guarantee. Since the performance bound of the algorithm given in the hint is difficult to prove, I designed another very similar algorithm that can easily prove its 3-approximation property.




Figure 3.1: An example of the scheduling algorithm

Algorithm Description:

Since $p = 2^k$ for some integer k , there are only $\log_2 p$ distinct scales of processor requirements. We partition jobs into categories based on their processor requirements:

- Category i : jobs with $P[j] \in (2^{i-1}, 2^i]$ for $i = 0, \dots, k$

The algorithm proceeds as follows:

1. categorize all jobs by processor requirements $P[j]$ into k categories
2. For each category i from k to 0:
 - Round up the processor requirements of all jobs in this category to 2^i
 - For each job in this category, find the processor with the earliest end time among all current processors, for example, processor j , and assign $p[j] \sim p[j + 2^i - 1]$ to this job. Such assignment would push the end time of these 2^i processors to the future.

Time Complexity:

Since there are in total n jobs, and each requires searching the earliest end time among p processors, which can be done in $O(\log p)$ time, the total time complexity is $O(n \log p)$.

3.3 Lower Bounds for Optimal Makespan

Let M^* be the optimal makespan. We establish two lower bounds:

Bound 1: $M^* \geq \max_{i=1}^n T[i]$

This is trivial: any job i must run for at least $T[i]$ time units. Therefore, the makespan cannot be smaller than the longest job duration.

Bound 2: $M^* \geq \frac{\sum_{i=1}^n T[i] \cdot P[i]}{p}$

This inequality represents the relationship between total computation capacity and job requirements, regardless of scheduling. The total computation capacity can be considered a large rectangle box, which is $p \times M^*$, need to fit in all jobs, each can be considered a small rectangle box, which is $P[i] \times T[i]$.

Thus, it is obvious that the volume of the large box is at least the sum of the volumes of all small boxes, that is, $p \times M^* \geq \sum_{i=1}^n T[i] \cdot P[i]$, which gives $M^* \geq \frac{\sum_{i=1}^n T[i] \cdot P[i]}{p}$.

3.4 3-Approximation Proof

We decompose the schedule produced by our algorithm into two parts:

$$\text{ALG} = A + B$$

where:

- A : total height of all non-last shelves across all categories
- B : total height of the last shelf in each category

$$B \leq \text{OPT}$$

Consider the last shelf in each category, which contains the largest duration job in that category,

Randomized Algorithm