

Dynamic Programming and Stochastic Control - Assignment 2

Zhenrui Zheng

Chinese University of Hong Kong, Shenzhen
225040512@link.cuhk.edu.cn

Contents

Exercise 1: Capacity Expansion over N Periods	1
1.1 Derivation of the DP Algorithm	1
1.2 (s, S) -type Optimal Policy	1
Exercise 2: Single-leg Revenue Management (Optimal Protection)	2
Exercise 3: EMSR-b Heuristic	3
Exercise 4: Parking Problem	4
4.1 Recursion	4
4.2 Threshold-Optimal Policy	4

Exercise 1: Capacity Expansion over N Periods

1.1 Derivation of the DP Algorithm

Let $J_k(x_k)$ be the optimal expected cost to go from period k to the terminal period N , given that the capacity at the beginning of period k is x_k . The Bellman equation is:

$$\begin{aligned} J_k(x_k) &= \min_{u_k, \dots, u_{N-1}} \mathbb{E} \left[-S(x_N) + \sum_{k=k}^{N-1} (C_k(u_k) + P_k(x_k + u_k - w_k)) \right] \\ &= \min_{u_k, \dots, u_{N-1}} \mathbb{E} \left[C_k(u_k) + P_k(x_k + u_k - w_k) - S(x_N) + \sum_{k=k+1}^{N-1} (C_k(u_k) + P_k(x_k + u_k - w_k)) \right] \\ &= \min_{u_k} \mathbb{E} \left[C_k(u_k) + P_k(x_k + u_k - w_k) + \min_{\dots, u_{N-1}} \mathbb{E} \left[-S(x_N) + \sum_{k=k+1}^{N-1} (C_k(u_k) + P_k(x_k + u_k - w_k)) \right] \right] \\ &= \min_{u_k} \mathbb{E} [C_k(u_k) + P_k(x_k + u_k - w_k) + J_{k+1}(x_k + u_k)] \end{aligned}$$

1.2 (s, S) -type Optimal Policy

From the Bellman equation in Section 1.1, we have:

$$J_k(x_k) = \min_{u_k \geq 0} \mathbb{E} [C_k(u_k) + P_k(x_k + u_k - w_k) + J_{k+1}(x_k + u_k)].$$

Let $y_k = x_k + u_k$ denote the capacity after expansion in period k . Then $u_k = y_k - x_k$ and the minimization becomes:

$$J_k(x_k) = \min_{y_k \geq x_k} [C_k(y_k - x_k) + \mathbb{E}[P_k(y_k - w_k)] + J_{k+1}(y_k)].$$

Define the function

$$G_k(y) = c_k y + \mathbb{E}[P_k(y - w_k)] + J_{k+1}(y),$$

where we note that $C_k(y - x_k) = K \mathbf{1}_{\{y > x_k\}} + c_k(y - x_k)$ for $y > x_k$, and $C_k(0) = 0$ for $y = x_k$.

Since S is concave with $\lim_{x \rightarrow \infty} dS(x)/dx = 0$, P_k are convex, and $c_k y + \mathbb{E}[P_k(y - w_k)] \rightarrow \infty$ as $|y| \rightarrow \infty$, the function $G_k(y)$ is convex and coercive. Therefore, $G_k(y)$ attains its global minimum at some point S_k^* .

The optimal policy is then:

- If $G(x_k) \geq G(S_k^*) + K$, expand to $y_k = S_k^*$ (i.e., $u_k = S_k^* - x_k$).
- If $G(x_k) < G(S_k^*) + K$, do not expand (i.e., $u_k = 0$).

This is precisely the (s, S) -type policy: when the cost of current capacity x_k falls below the threshold $G(S_k^*) + K$, maintain the current capacity; otherwise, expand to the target level S_k^* .

Exercise 2: Single-leg Revenue Management (Optimal Protection)

Omitting the derivation, we provide that the optimal control at stage $j + 1$ is:

$$\begin{aligned}\mu_{j+1}^*(x, D_{j+1}) &= \min \{(x - y_j^*)^+, D_{j+1}\} \\ y_j^* &= \max \{x : 0 \leq x \leq C, p_{j+1} \leq \Delta V_j(x)\}\end{aligned}$$

The corresponding cost function is:

$$V_j(x) = \mathbb{E} [p_j \min \{(x - y_{j-1}^*)^+, D_j\} + V_{j-1}(x - \min \{(x - y_{j-1}^*)^+, D_j\})]$$

Which can be implemented as follows:

1. Iterate over stage j from 1 to 10 (Backward).
2. For each stage j , compute the optimal level y_j^* with $O(1)$ amortized complexity.
3. Iterate over capacity x from 0 to C .
4. Compute the expected revenue $V_j(x)$, which is $O(C)$ complexity.

The result of optimal protection levels and total expected revenue are:

$$V_{10}(100) \simeq 35415.63, \quad y_j^* = \{6, 15, 26, 36, 47, 57, 69, 78, 91\}$$

An example of python implementation is provided below:

```
import numpy as np
from scipy.stats import norm

# Normalized discretized truncated normal distribution
mu, sigma = 10, 2
p = np.zeros(21, dtype=np.float64)
for k in range(21):
    p[k] = norm.cdf((k+0.5-mu)/sigma) - norm.cdf((k-0.5-mu)/sigma)
p = p / np.sum(p)

C = 100
prices = np.array([0, 500, 480, 465, 420, 400, 350, 320, 270, 250, 200], dtype=np.float64)

V = np.zeros((11, C+1), dtype=np.float64)
y_max = 0
def binary_search(l, r, j):
    # find the largest x s.t. V[j, x] - V[j, x-1] >= prices[j+1]
    while l < r-1:
        m = (l + r) // 2
        if V[j, m]-V[j, m-1] >= prices[j+1]:
            l = m
        else:
            r = m
    return l

for j in range(1, 11):
    y_max = binary_search(y_max, C, j-1)
    print(f"y{j-1} = {y_max}")
    for x in range(C+1):
        for D in range(21):
            V[j, x] += p[D] * (prices[j] * min(max(x-y_max, 0), D) + V[j-1, x-min(max(x-y_max, 0), D)])
print(f"V[10, 100] = {V[10, 100]}")
```

Exercise 3: EMSR-b Heuristic

Following the description in the problem, we can compute the optimal protection levels and total expected revenue in a modified version:

1. Iterate over stage j from 1 to 10 (Backward).
2. For each stage j , compute the aggregated future demand $\sum_{k=1}^{j-1} \mathbb{E}[D_k]$ and weighted-average revenue $\bar{p}_{j-1} = \frac{\sum_{k=1}^{j-1} p_k \mathbb{E}[D_k]}{\sum_{k=1}^{j-1} \mathbb{E}[D_k]}$.
3. Since the demand is i.i.d. normal, $\mu_S = (j - 1) \times \mu$ and $\sigma_S^2 = (j - 1) \times \sigma^2$.
4. Compute the constant $z_\alpha = \Phi^{-1}(1 - p_j / \bar{p}_{j-1})$ and $y_j = \mu_S + z_\alpha \sigma_S$.
5. Iterate over capacity x from 0 to C .
6. Compute the expected revenue $V_j(x)$.

The result of optimal protection levels and total expected revenue are:

$$V_{10}(100) \simeq 35413.23, \quad y_j^* = \{6, 15, 26, 36, 47, 57, 68, 78, 90\} \quad (\text{rounded to nearest integer})$$

An example of python implementation is provided below:

```
# ESMR-b
mu_sum = 0
sigma_sum = 0
total_revenue = 0
V = np.zeros((11, C+1), dtype=np.float64)
y_max = 0
for j in range(1, 11):
    if j > 1:
        mu_sum += mu
        sigma_sum += sigma**2
    total_revenue += prices[j-1] * mu
    p_bar = total_revenue / mu_sum
    z_alpha = norm.ppf(1-prices[j]/p_bar)
    y_max = mu_sum + z_alpha * sigma_sum**0.5
    y_max = math.ceil(y_max-0.5)
    print(f"y{j-1} = {y_max}")
    for x in range(C+1):
        for D in range(21):
            V[j, x] += p[D] * (prices[j] * min(max(x-y_max, 0), D) + V[j-1, x-min(max(x-y_max, 0), D)])
print(f"V[10, 100] = {V[10, 100]}")
```

Exercise 4: Parking Problem

Each parking spot is independently free with probability p . The cost of parking k spots away from the destination is k ; reaching the destination without parking incurs cost C . Let F_k be the minimal expected cost when k spots from the destination, with $F_0 = C$.

4.1 Recursion

At distance k , if the current spot is free (probability p), the driver may

- park now, incurring cost k ; or
- pass it and continue, incurring expected cost F_{k-1} .

If the spot is occupied (probability $q = 1 - p$), the driver must continue with cost F_{k-1} . Hence

$$F_k = p \min\{k, F_{k-1}\} + q F_{k-1}, \quad k = 1, 2, \dots$$

and $F_0 = C$.

4.2 Threshold-Optimal Policy

From the recursion, when $k \leq F_{k-1}$ it is better to park; otherwise it is better to continue. Therefore there exists a threshold k^* such that

if $k < k^*$ then park at the first free spot; if $k \geq k^*$ then never park.

By monotonicity and $F_0 = C$, we obtain

$$k^* = \min \left\{ i \in \mathbb{Z}_{\geq 0} : q^i < \frac{1}{pC + q} \right\},$$

which is equivalent to the expression given in the exercise (note $q = 1 - p$). Thus the threshold parking policy is optimal.